


Publisher Team

Mission: find & buy 300+ million profitable ad placements every day

- Cookie Targeting: the work horse
 - 100 billion publisher arbitrations every day
 - 5 billion publisher pixels called every day
- Real-Time Bidding: the race horse
 - Started from scratch in May'10
 - Integrated since then with 8 RTB partners
 - 01/09/2011: 1.95 billion requests, 1.55 billion bids
- Criteo Publisher Marketplace : the mustang
 - (Re)started from scratch in Oct'10
 - Self-service for 2500+ publishers in 6 countries

1.9 million RTB requests per minute



Publisher Team: what's next

- Cookie Targeting
 - Quite mature, no huge new features on the radar
 - Major upgrades under the hood (Scalability team, more in a minute)
- Real-Time Audience: the next step in Cookie Targeting
 - Better arbitrage accuracy, pixel-free solution
 - Free retargeting alternative for publishers (DoubleClick)
 - Live on 16 publishers (inc. SevenOne, #3 in DE)



Publisher Team: what's next

- Criteo Publisher Marketplace
 - A zillion new features (no less)
 - Productivity features for Criteo teams
- Real-Time Bidding
 - More RTB partners (Yahoo coming)
 - Ramp-up in Europe, launch in Japan
 - Inventory quality : ad verification, blacklisting
 - Productivity tools for Criteo teams : publisher clustering, better stats, smarter logs, new Dashboard, etc.

Publisher Team: what's next



- Spin-off a new team in charge of hunting data
 - Use data already available in RTB requests (Google keywords, geotargeting, ...)
 - Import socio-demographics (Quantcast)
 - Crawl publisher pages to extract valuable data (page content) and metadata (# of banners, above/below the fold, etc.)
 - Extract valuable keywords from URLs
 - Extract valuable keywords from Advertiser catalogs
- All this could be used to optimize buying
- Several experiments already in progress

Scalability Team

Mission: if it ain't broken, fix it before it breaks. If it IS broken, fix it already!


Main lines of work:

- Develop critical software modules used by all R&D teams
- Bleeding-edge technologies, from PoC to production
- Provide platform-wide coaching & expertise on performance topics
- Code refactoring
 - 50% drop in memory consumption on Cookie targeting servers (left)
 - 80% drop in memory consumption on Delivery servers (right)



Scalability Team

- Catalog Import v2
 - Able to ingest & replicate monster US catalogs coast-to-coast
 - Completely redesigned
- Real-time publisher arbitration
 - Replaces pre-computed files loaded on web servers
 - Less latency in tagging/detagging
- Image delivery to banners: 690 million daily
- Log aggregation
 - over 1TB crunched daily
 - Hadoop cluster (700 cores, 700 TB)



Name	Size	User	Group	Permissions	Date
hadoop	0 bytes	rtb	hdfs	drwxr-xr-x	Aug 27, 2011 1:15 a.m.
KeywordsArchives	0 bytes	keywords	hdfs	drwxr-xr-x	Sept 5, 2011 7:45 p.m.
KeywordsBestOfs	0 bytes	keywords	hdfs	drwxr-xr-x	Aug 30, 2011 7:19 p.m.
KeywordsMetrics	0 bytes	keywords	hdfs	drwxr-xr-x	Sept 5, 2011 9:09 a.m.
KeywordsRatings	0 bytes	keywords	hdfs	drwxr-xr-x	Sept 2, 2011 9:12 p.m.
benchmarks	0 bytes	hdfs	hdfs	drwxr-xr-x	June 7, 2011 9:51 p.m.
current_run	0 bytes	rtb	hdfs	drwxr-xr-x	Aug 27, 2011 1:15 a.m.
hadoop	0 bytes	rtb	hdfs	drwxr-xr-x	Aug 26, 2011 7:24 p.m.
hdfs	0 bytes	rtb	hdfs	drwxr-xr-x	July 11, 2011 8:52 a.m.
home	0 bytes	rtb	hdfs	drwxr-xr-x	Sept 2, 2011 12:27 p.m.
last_run	0 bytes	rtb	hdfs	drwxr-xr-x	Aug 26, 2011 1:15 a.m.
system	0 bytes	rtb	hdfs	drwxr-xr-x	Sept 1, 2011 4:55 p.m.

Name / Id	Status	User	Maps/Reducers	Queue	Priority	Duration	Date
file-import-kawari	Running	file	100/100	default	normal	16m 1s	09/05/11 15:12:42

Production team

- Mission: 100% uptime, 24/7/365
- Design, procurement & setup:
 - Datacenters (5)
 - Servers + HPC (830)
 - Networking equipment & links
 - Databases (close to 4000)
- Monitoring and on-duty calls active on all of the above, as well as applications



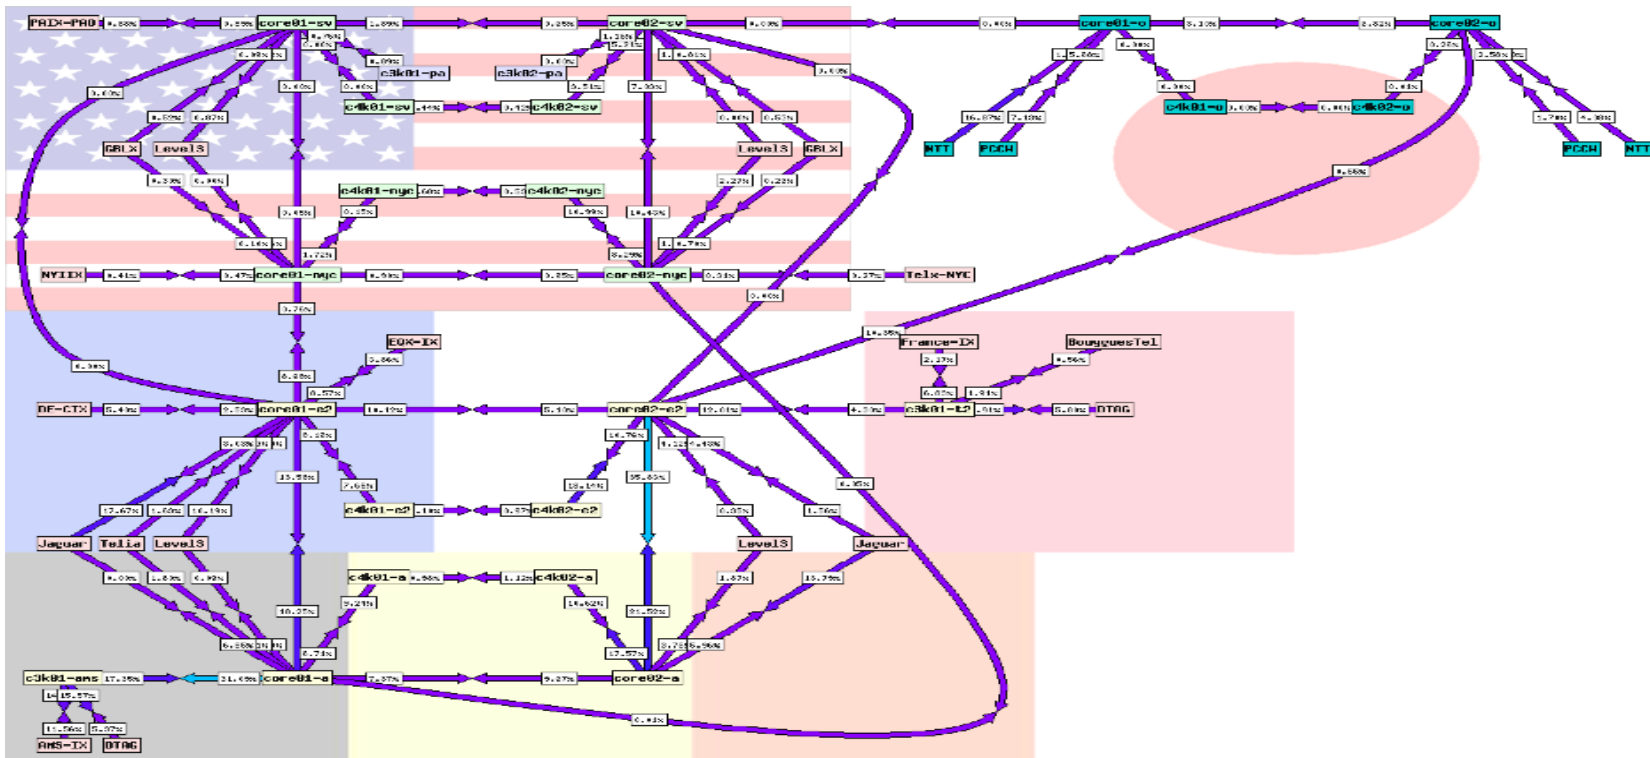
Production team: hosting

New site outside Paris (40 racks)
This one's our blueprint






Production team: network

- Up to **4Gb/s** of traffic in Europe
- Up to **1Gb/s** of traffic in the USA
- Up to **300Mb/s** of traffic in Asia
- 200 peering sessions



Created: Sep 20 2011 11:00:00

Production team: traffic



Peering

- Some references among hundreds of peers!



swisscom

TELE2



Production Team: what's next

- 2011
 - Finish migration & cleanup of Paris DC
 - Scale Tokyo DC (already?)
 - Double our hosting capacity in the US (NY & SV)
- 2012
 - Migrate Brussels DC to Amsterdam
 - Deploy HPC cluster in the US
 - 2nd datacenter in Asia
 - 3rd datacenter in the US (RTB-specific maybe?)
 - Eastern Europe, South America, ... we'll see ☺