

O'REILLY®

# Velocity CONFERENCE

BUILD RESILIENT SYSTEMS AT SCALE

[velocityconf.com/eu](http://velocityconf.com/eu)

#velocityconf

## A real-life account of moving 100% to a public cloud

Antoine Guy, Principal Infrastructure Architect  
Julien Simon, Chief Technology Officer



# A word about us

## Julien Simon

Last 10 years as VP Eng / CTO  
in web startups (DigiPlug, Pixmania,  
Criteo, Viadeo)

Disclaimer : I joined AWS two weeks ago,  
but don't worry, this won't be a sales pitch :D

Email: [julsimon@amazon.com](mailto:julsimon@amazon.com)

Twitter: @julsimon



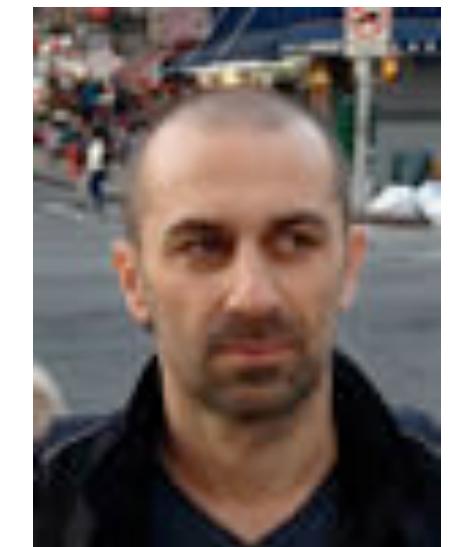
## Antoine Guy

14 years of infrastructure & web ops

Tech lead for Infrastructure and  
Operations at Viadeo since 2011

Email: [aguy@viadeoteam.com](mailto:aguy@viadeoteam.com)

Twitter: @anto1ne



## Viadeo.com

Leading professional social network  
Founded in 2005, 10M members in France

# What this case study will cover

- The rationale that led to an all-in move to Amazon Web Services
- The design options and trade-offs
- Why we did or didn't alter parts of our stack
- The complete automation of our infrastructure using AWS CloudFormation and continuous integration

# Technology stack

- Service-oriented platform: « Kasper » (Java)
  - CQRS: Command-Query Responsibility Segregation
  - ES: Event Sourcing
  - DDD: Domain Driven Design
- Web application: « Limbo » (node.js)
- Mobile applications: iOS and Android
- Backends: MySQL, Elasticsearch, Hbase, Spark, Hadoop
- Infrastructure: 15 racks hosted in San Francisco, 250 servers, limited use of virtualization



# Infrastructure challenges

1. Aging, unsupported hardware
2. Discrepancy between software agility and infrastructure rigidity
3. Operating a physical network
4. Operating a remote data center
5. Limited disaster recovery capabilities ( $RTO > 1$  day)

# 4 options

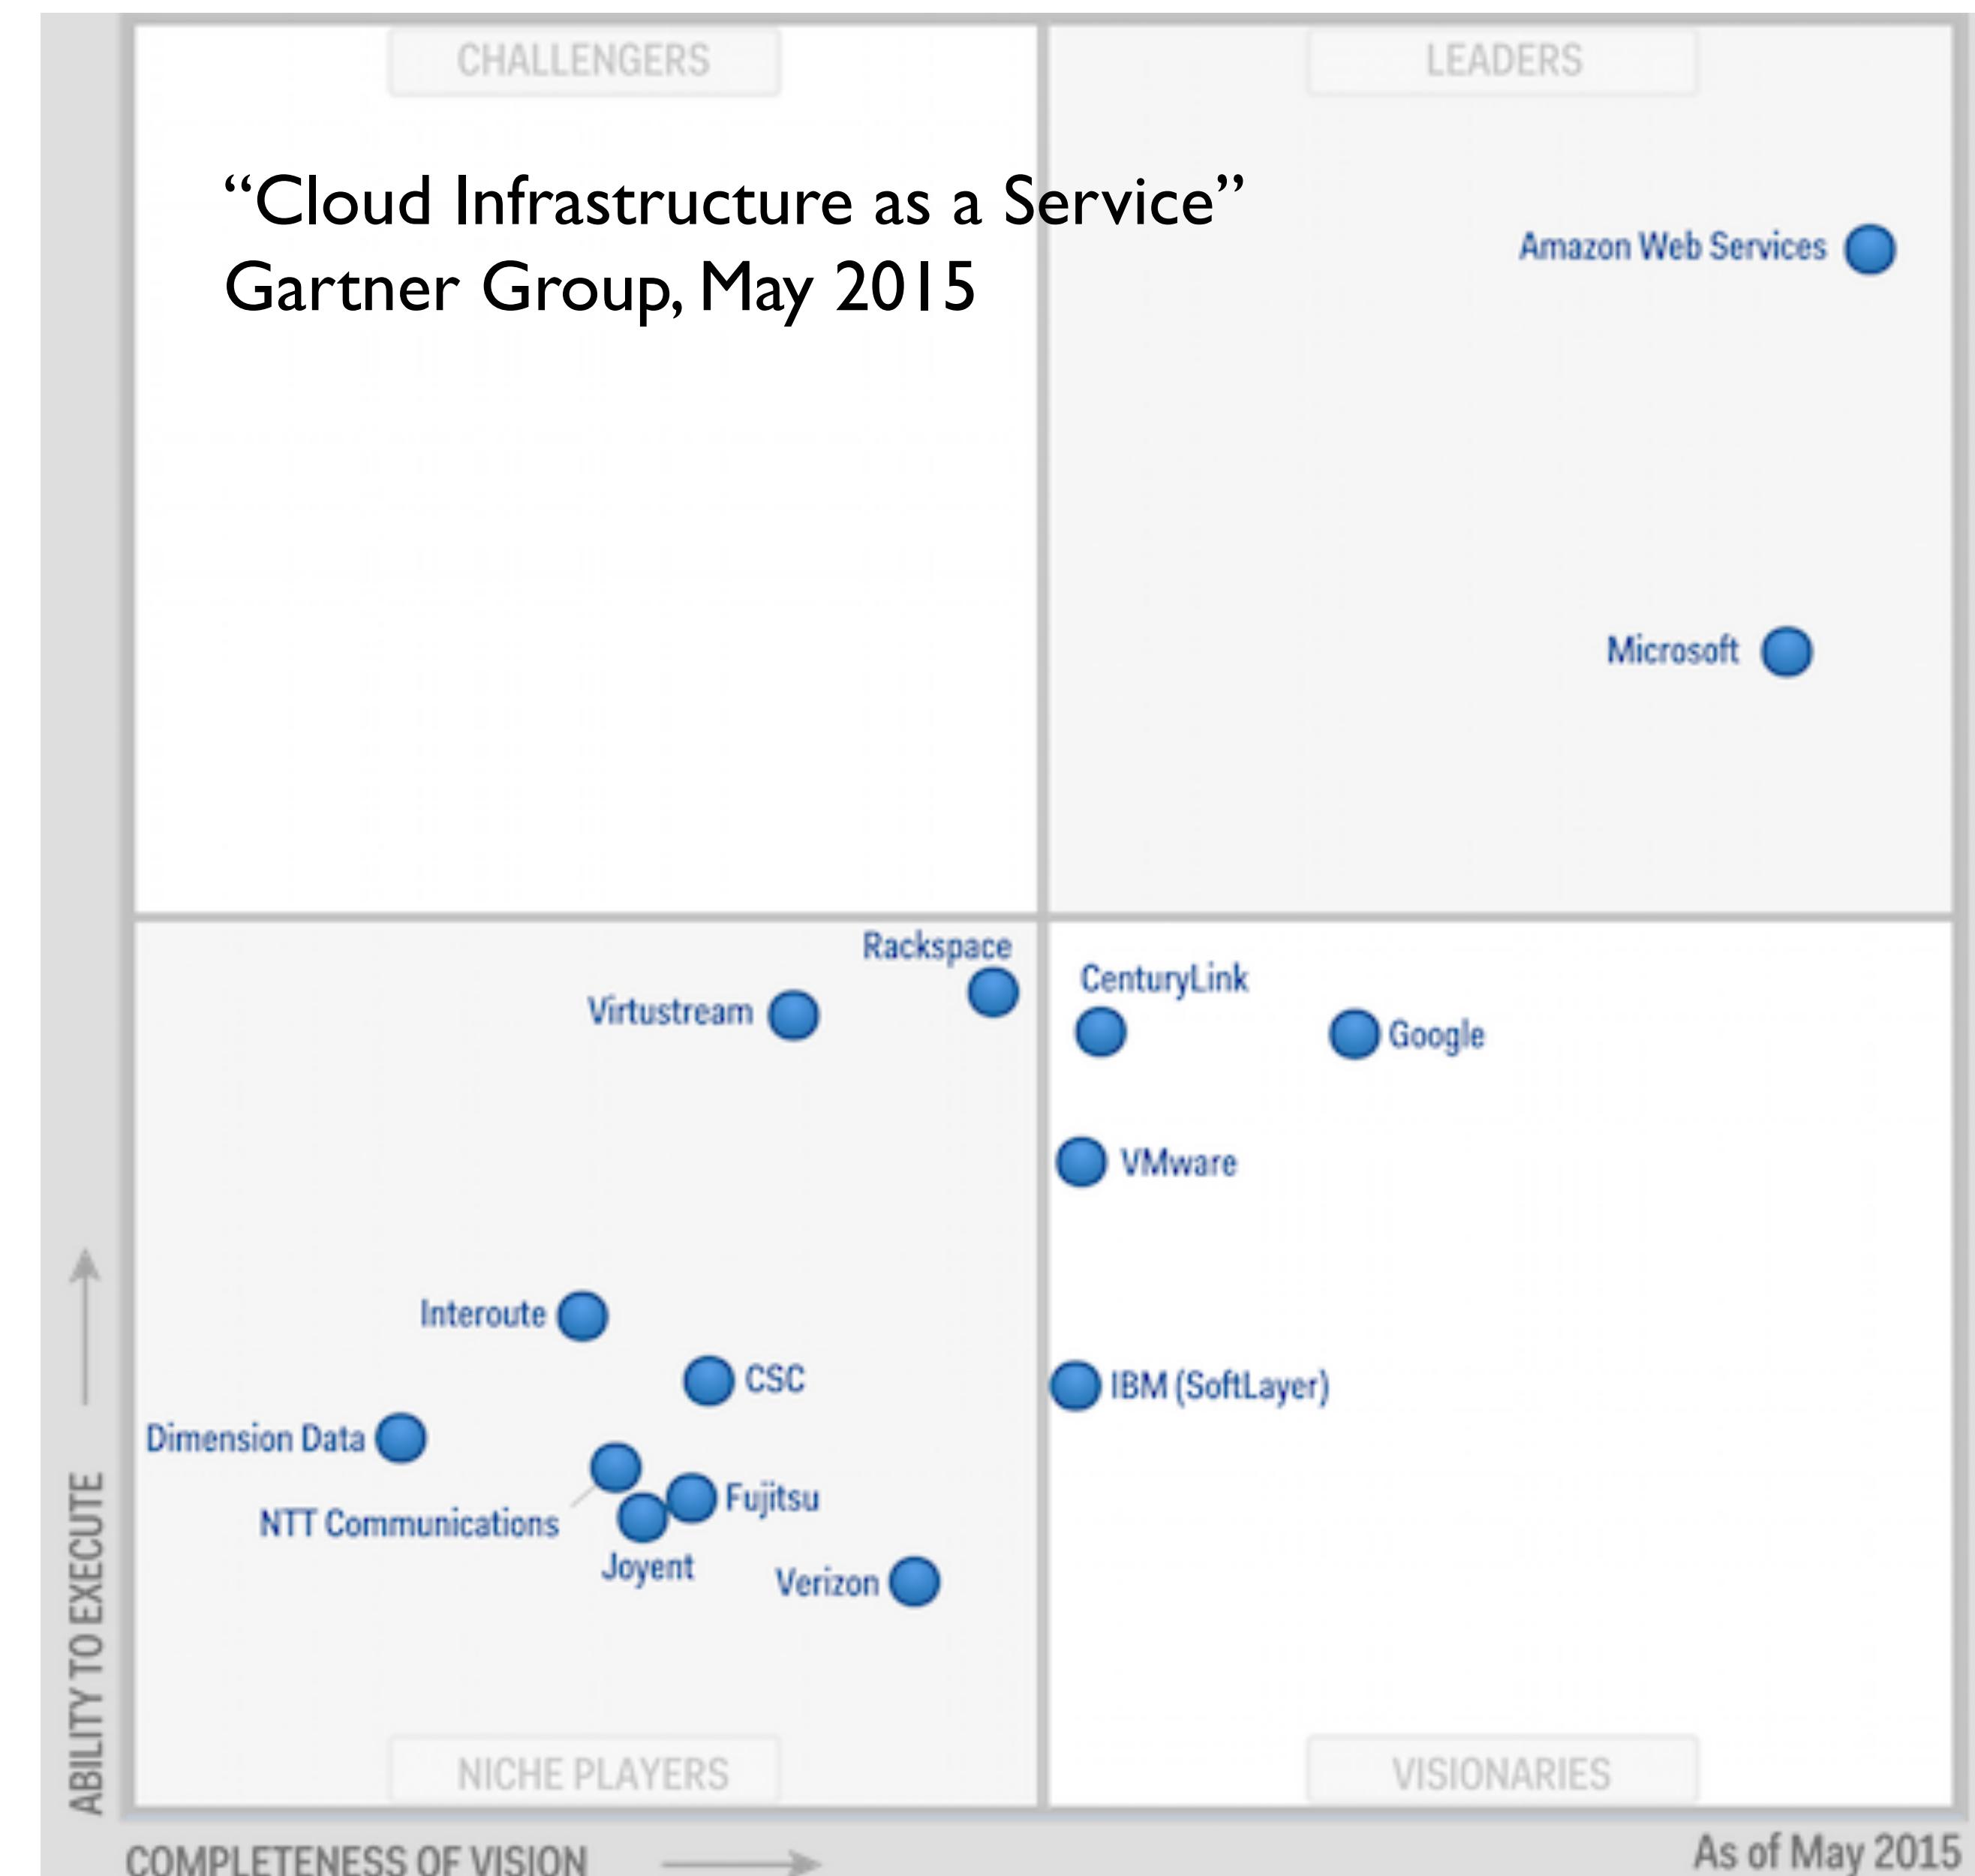
|                  | Server upgrade | Full virtualization | New datacenter | Public cloud  |
|------------------|----------------|---------------------|----------------|---------------|
| Refresh hardware | Solved         | Solved              | Solved         | Solved        |
| Improve agility  | No change      | Better              | No change      | Better        |
| Manage network   | No change      | Worse               | No change      | Better        |
| Manage remote DC | No change      | No change           | No change      | Gone          |
| Improve DR       | No change      | Little better       | Better         | Solved        |
| Effort needed    | Low            | Medium              | High           | Damn!         |
| Learning curve   | None           | Low                 | Low            | Ouch!         |
| Build spend      | Very high      | High                | Very high      | None          |
| Run spend        | Little better  | Better              | Better         | Pay as you go |
| Gut feeling      | Meh...         | Too late            | Wasteful       | Party on!     |

# Why a public cloud?

- We want to focus on our real job : **building a great service**
  - No more issues with hardware vendors, no more licensing hell. Ever!
  - Experimenting, deploying & scaling are just a few clicks away
- We're **ready and eager for it** : Agile and DevOps are Viadeo values
- We want to **measure ROI and optimize cost**
- We've been using it for a while now **and it just works**

# Why AWS

- Already using « cold » services
- Powerful PaaS technology accessible in a few clicks
- Multi region
- Lots of new features, constantly improving



# Play book

- Key objectives: **automation, scalability, safety**
- Continuous integration & delivery **everywhere** (infra, config, apps)
- Replace parts of our stack only if the **benefits** are too good to pass
- All-in, but **no fork lift**: move gradually (and be ready to rollback)
- Plan & build with **temporary hybrid** run in mind

# Pre-flight check

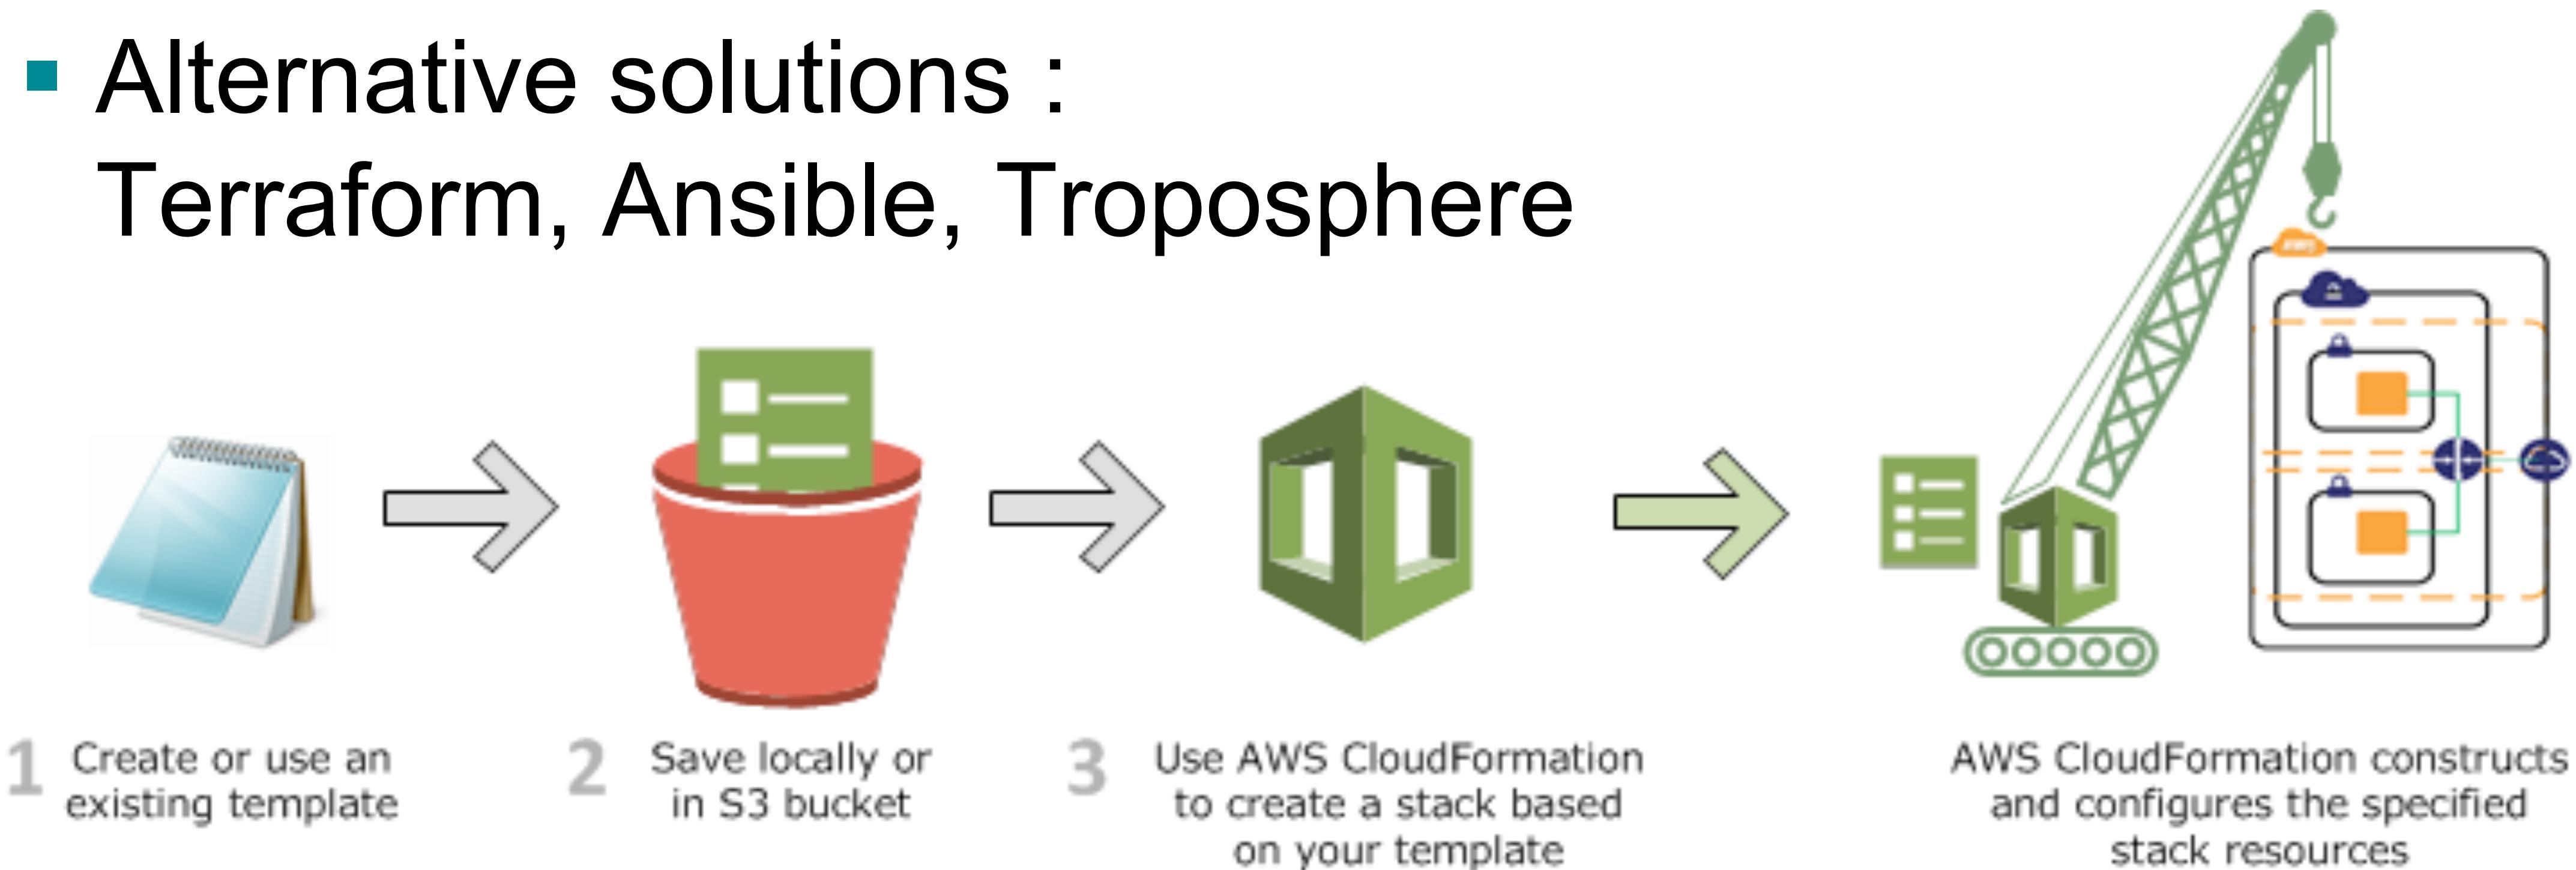
- Build a **thorough** report of your infrastructure.
  - Estimate equivalent cost in cloud. Don't be afraid by initial amount.
  - Evaluate each part for replacement: leave as is, PaaS, SaaS ?
  - Find pain points: technical debt, relevance of moving legacy apps.
- Define a **high-level** migration plan (that you won't follow)
  - Usually : staging, analytics, new projects, critical/legacy, close DC.
  - Expect (and allow) early adopters or emergencies to go first.

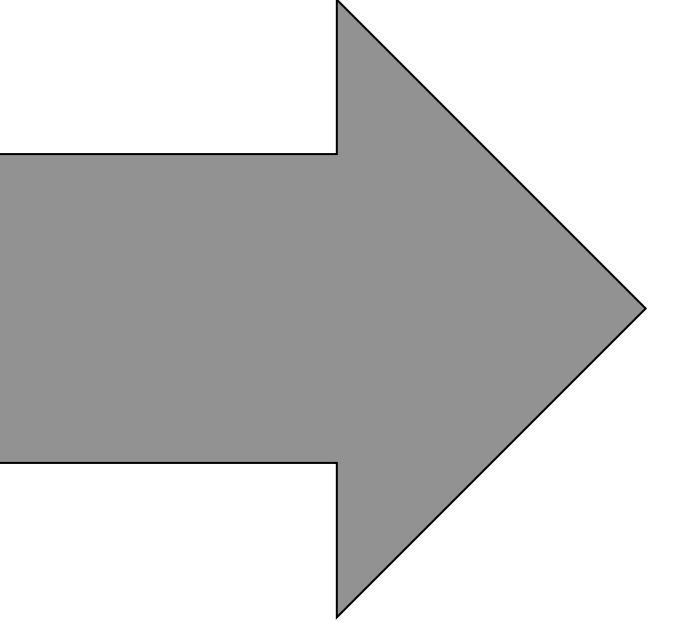
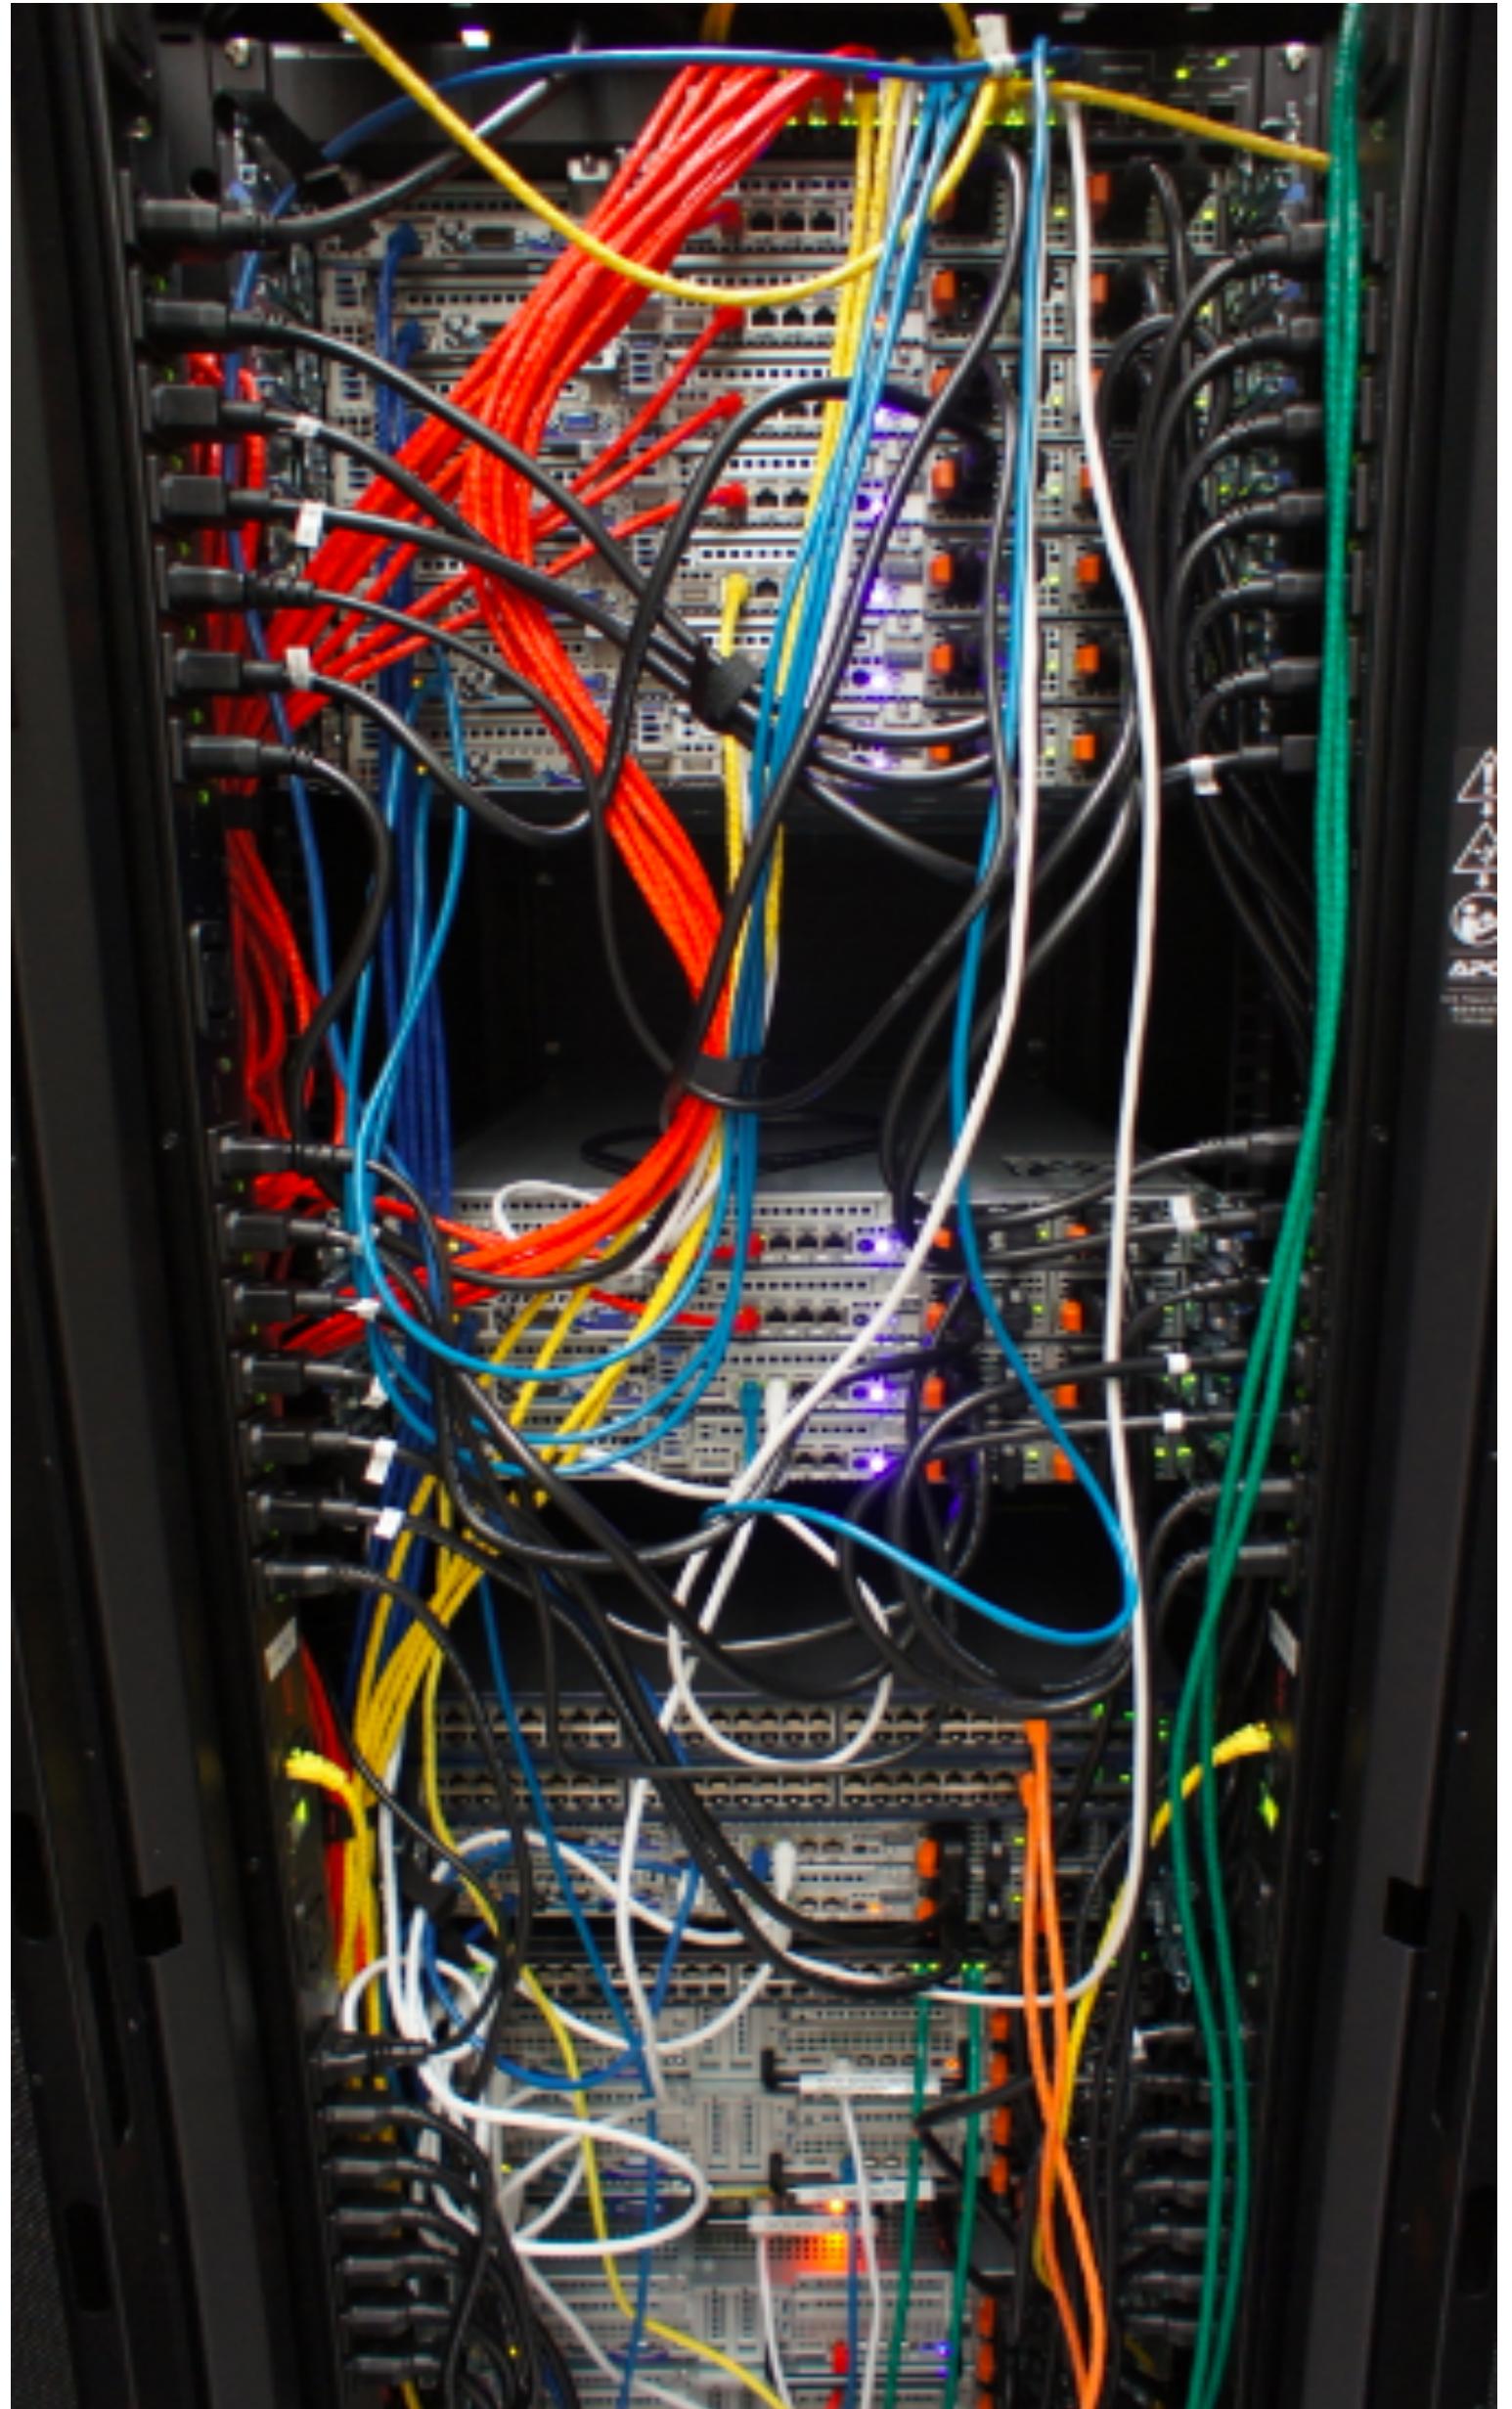
# Infrastructure as code, for real

- Versioned, auditable blueprints (developers can contribute)
- Quick to deploy, repeatable, tested infrastructure
- Enables CI/CD for infrastructure (just like everything else)
- Deploy anywhere (for multi-region cloud)
- Cloud can now be its own DR

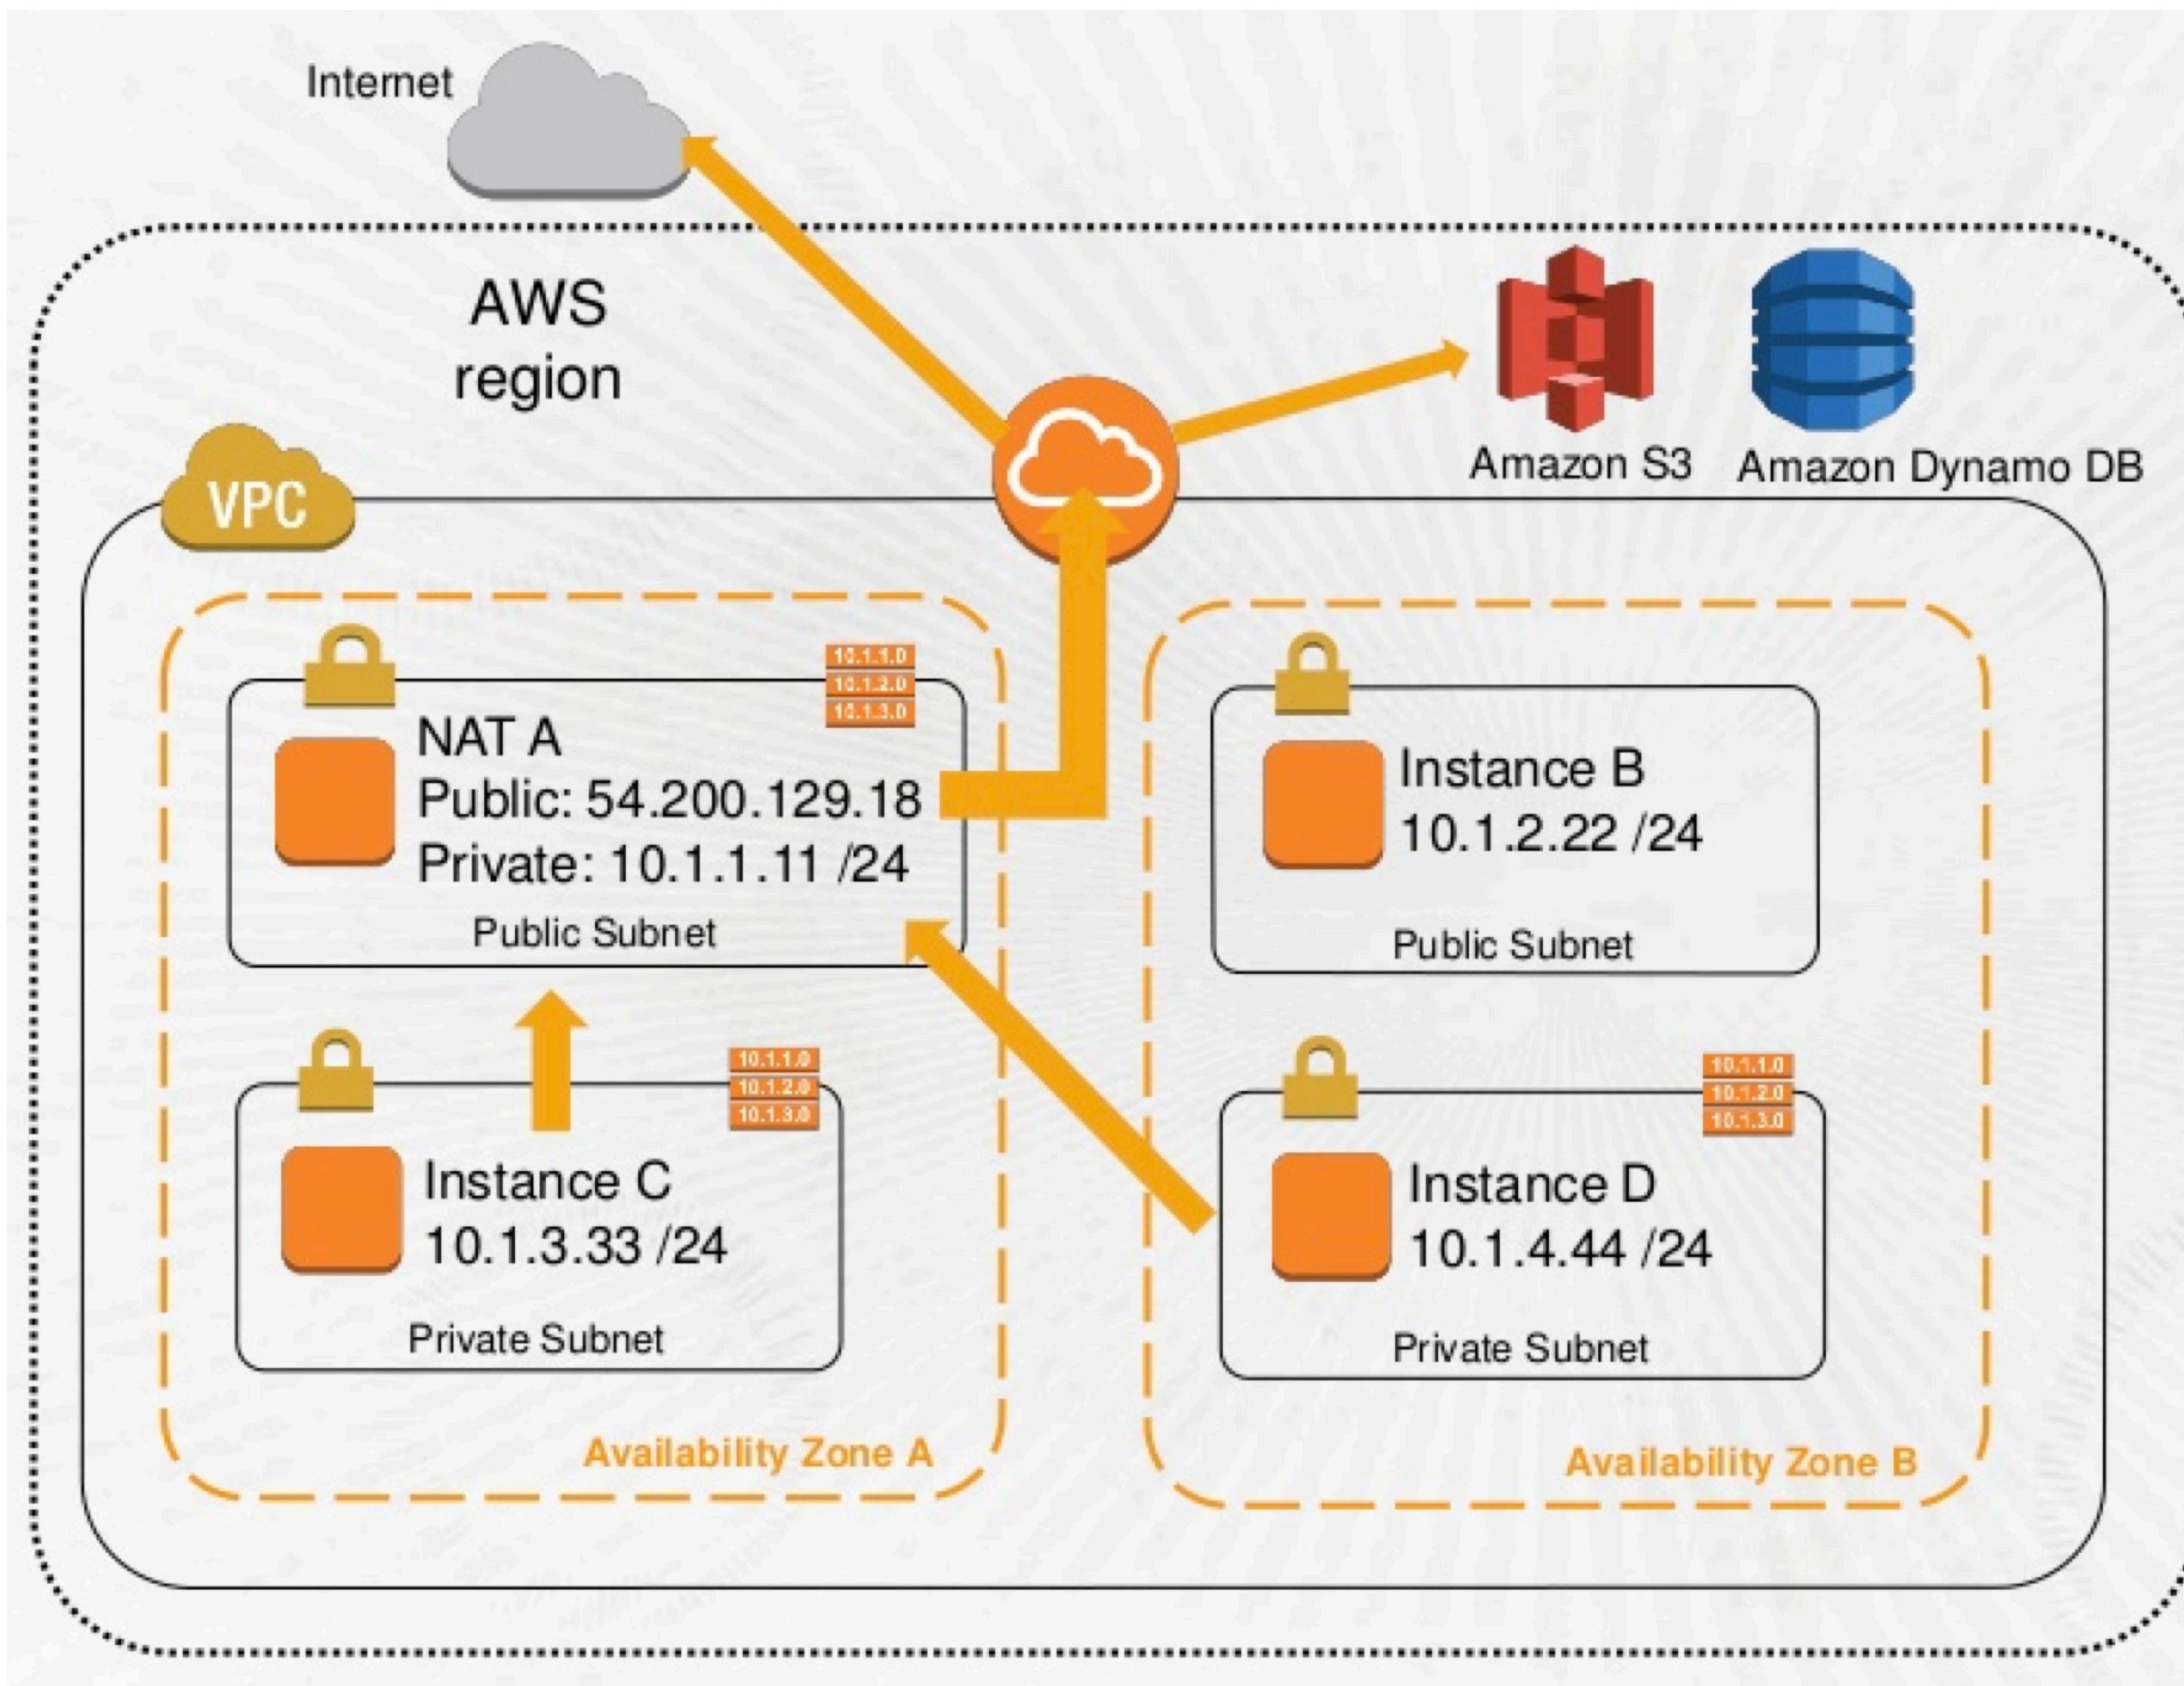
# AWS CloudFormation

- **Json-based** descriptive “language” to create a set of AWS resources called a **stack**.
- Alternative solutions :  
Terraform, Ansible, Troposphere





```
"Conditions" : {
    "HaveNoOtherRoles" : { "Fn::Equals" : [{"Ref" : "OtherRoles"}, """] },
    "HaveEbs" : { "Fn::Not" : [{ "Fn::Equals" : [{"Ref" : "EbsVolumeSize"}, "0"]} ] },
    "HaveEbsSnapshotId" : { "Fn::Not" : [{ "Fn::Equals" : [{"Ref" : "EbsSnapshotId"}, """]} ] },
    "HaveAdditionalTagKey": { "Fn::Not" : [{ "Fn::Equals" : [{"Ref" : "AdditionalTagKey"}, """]} ] },
    "HaveAdditionalTagValue": { "Fn::Not" : [{ "Fn::Equals" : [{"Ref" : "AdditionalTagValue"}, """]} ] },
    "HaveSSL": { "Fn::Not" : [{ "Fn::Equals" : [{"Ref" : "SSLPort"}, "0"]} ] },
    "IsHTTP" : { "Fn::Equals" : [{"Ref" : "ElbProtocol"}, "HTTP"] },
    "HaveSpotPrice" : { "Fn::Not" : [{ "Fn::Equals" : [{"Ref" : "SpotPrice"}, """]} ] }
},
"Resources": {
    "AutoScalingGroup": {
        "Type": "AWS::AutoScaling::AutoScalingGroup",
        "UpdatePolicy" : {
            "AutoScalingRollingUpdate" : {
                "MaxBatchSize" : "1",
                "MinInstancesInService" : "0",
                "PauseTime" : "PT15M",
                "WaitOnResourceSignals": "true"
            }
        },
        "Properties": {
            "LaunchConfigurationName": { "Ref": "LaunchConfig" },
            "LoadBalancerNames": [ { "Ref": "ElasticLoadBalancer" } ],
            "MinSize": { "Ref": "MinPoolSize" },
            "MaxSize": { "Ref": "MaxPoolSize" },
            "AvailabilityZones": { "Fn::FindInMap": ["AZConfig", "AvailabilityZones", "all"] },
            "VPCZoneIdentifier": { "Ref": "EC2SubnetsIds" },
            "Tags" : [
                { "Fn::If": [
                    "HaveAdditionalTagKey",
                    {
                        "Key" : { "Ref": "AdditionalTagKey" },
                        "Value": {
                            "Fn::If": [
                                "HaveAdditionalTagValue",
                                {"Ref": "AdditionalTagValue"},
                                ""
                            ]
                        },
                        "PropagateAtLaunch": "true"
                    },
                    {"Ref" : "AWS::NoValue"}
                ]
            },
            { "Key" : "Name", "Value" : { "Fn::Join" : [ ".", [ { "Ref" : "ServiceName"}, {"Ref" : "EnvironmentName"} ] ] },
            { "Key" : "cost", "Value" : { "Ref" : "Cost" }, "PropagateAtLaunch": "true" },
            { "Key" : "environment", "Value": { "Ref" : "EnvironmentName"}, "PropagateAtLaunch": "true" }
        ]
    }
}
```

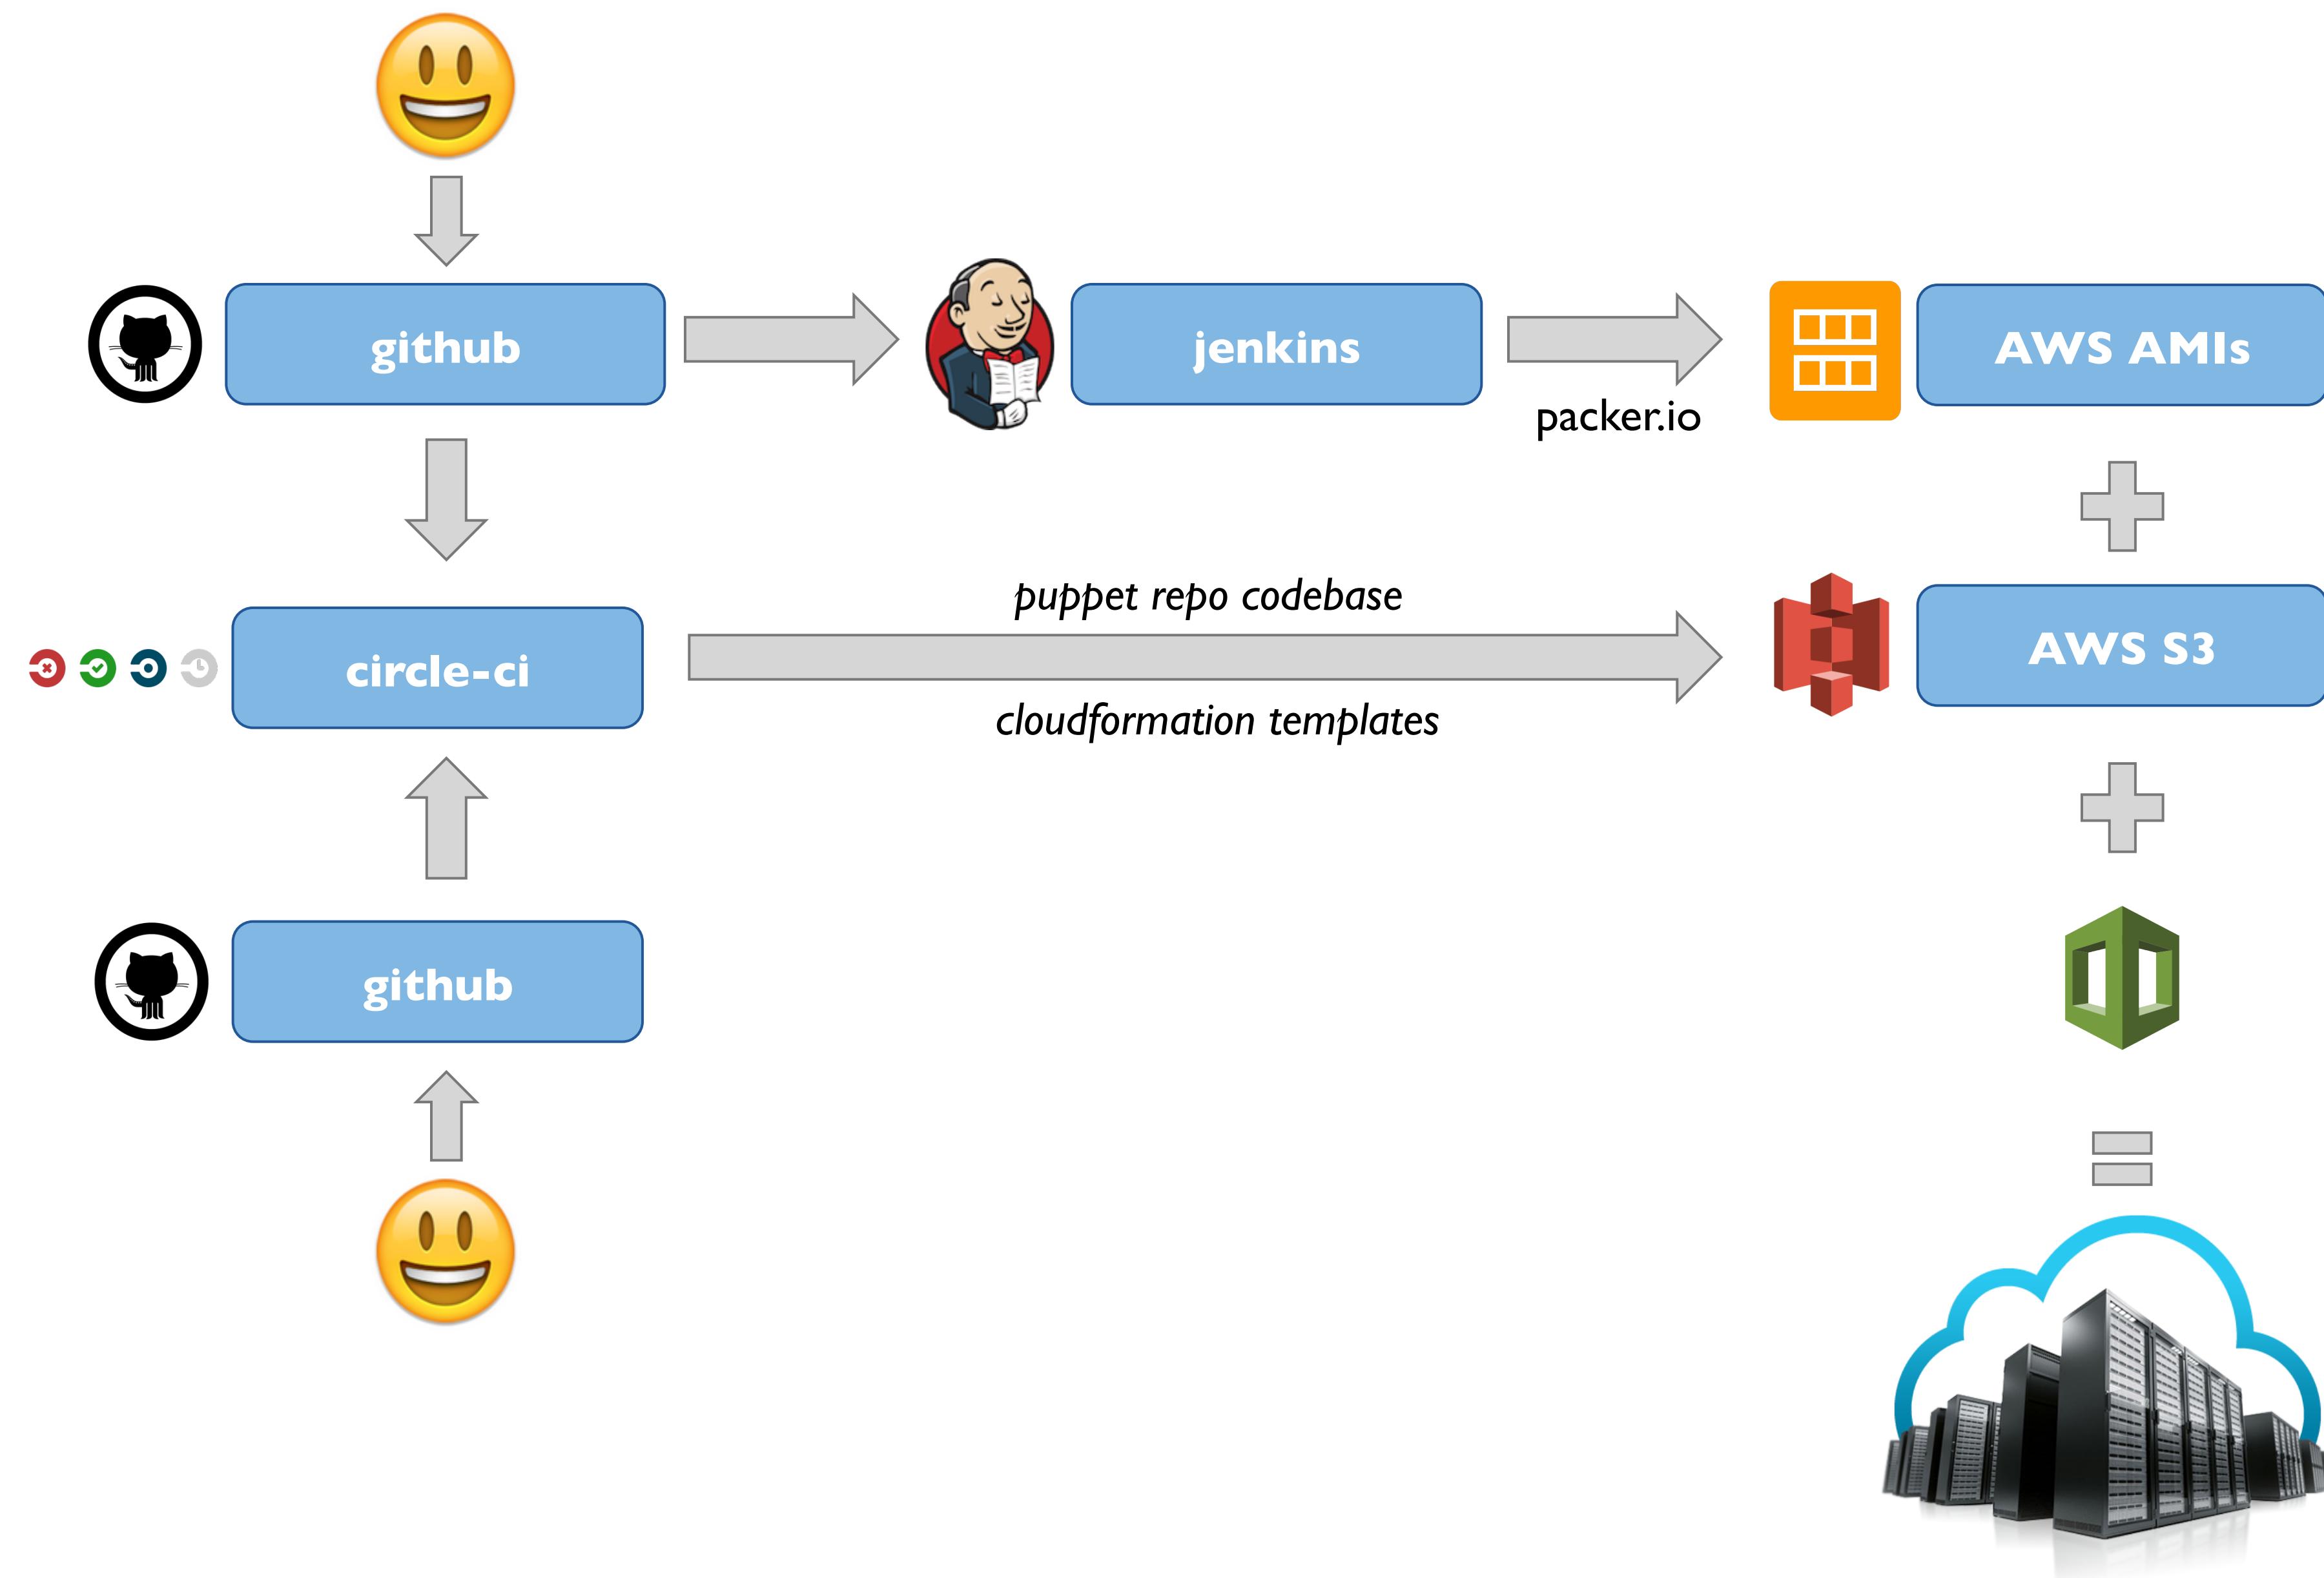


# Step 1: automate your infra build

- `$ aws cloudformation validate-template`
- Write your test cases in the language of your choice (many SDKs available)
- Use git with a trigger to a CI tool (`circle-ci` or other)
- If tests pass : deploy (in your staging environment first)

# Step 2: build and CI your images

- **Automate** your baking:  
1 for each role from your config management
- **Bake** automatically with packer or aminator
- **Baked or bootstrapped ?**
  - Fully baked for autoscale services
  - Bootstrapped for non critical, fast-changing services



# CloudFormation best practices

- Use AWS CloudFormer as a **starting point**
- **Reuse** as much as possible
  - Nested stack, but not too much, max 1 level (`hello UPDATE_ROLLBACK_FAILED`).
  - Parameterized template : env, region, name, size, etc..
- Use stacks to perform **green/blue** deployments
- Tag everything!

# AWS: lessons learned (1)

- Plan your VPC network **early**. Watch ARC403
- Hybrid run requires good **connectivity**
  - DirectConnect helps, but watch NET406
- ELB is **level 3** only, with limited options
- CloudFront performance not great across continents
- A 2<sup>nd</sup> CDN may be required, if only for peace of mind

# AWS: lessons learned (2)

- EMR != 24/7 Hadoop cluster
- EMR in VPC with private DNS is very difficult
- MySQL: RDS doesn't allow myISAM
- Limits: **check** them, raise them. Same everywhere
- No more static naming/addressing: use **service discovery**
- Make sure everything is **multi-AZ** and **tagged**

# Tech is only half the work

- Identify high-level **stakeholders** and understand **THEIR** goals
- Involve your **Legal/Finance** departments **early**
  - Budgeting: cost overlap, AWS support, etc.
  - Early termination of legacy infrastructure contracts (there will be blood)
- Build an **A-team** (dev, ops, security)
- Work on **awareness** and **knowledge transfer**

# Current status

- **Staging** environment: done and running
- **Production** web traffic served from AWS
- Infrastructure live in 3 regions (us-east-1, us-west-2, eu-west-1)
- Multiple inter-connected VPC
- ≈ 100 AWS EC2 instances (half production, half development)
- 2 AWS Redshift clusters (5 instances each)

# Next steps



- Start moving backend servers to AWS
- Live Hadoop cluster → EMR
- Clean-up MySQL stack before we can use RDS
- Elasticsearch will remain EC2-based for now
- Optimize: reserved instances, spot instances, auto-scale

# Conclusion

- 5 years ago : «Cloud computing? Why?»
- Now : «NO cloud computing? Why?»
- Some good reasons, but many outdated ones
- Cloud computing a fashion? We don't think so
- Infrastructure is now digital, like photos, music, movies, money, etc.
- We're making the most of it, it's a great engineering challenge!

O'REILLY®

# Velocity

CONFERENCE

BUILD RESILIENT SYSTEMS AT SCALE

aguy@viadeoteam.com

julsimon@amazon.com

[velocityconf.com/eu](http://velocityconf.com/eu)

#velocityconf

